

# Cars Price Prediction Regression Model

## ▼ Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
sns.set_style('whitegrid')
plt.style.use("fivethirtyeight")
```

```
data = pd.read_csv('cars_price.csv')
```

```
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	

```
data.shape
```

```
(205, 26)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	symboling	205 non-null	int64
1	normalized-losses	205 non-null	object
2	make	205 non-null	object
3	fuel-type	205 non-null	object
4	aspiration	205 non-null	object
5	num-of-doors	205 non-null	object
6	body-style	205 non-null	object
7	drive-wheels	205 non-null	object
8	engine-location	205 non-null	object
9	wheel-base	205 non-null	float64
10	length	205 non-null	float64
11	width	205 non-null	float64
12	height	205 non-null	float64
13	curb-weight	205 non-null	int64
14	engine-type	205 non-null	object
15	num-of-cylinders	205 non-null	object
16	engine-size	205 non-null	int64
17	fuel-system	205 non-null	object

```

18 bore                205 non-null object
19 stroke              205 non-null object
20 compression-ratio   205 non-null float64
21 horsepower          205 non-null object
22 peak-rpm            205 non-null object
23 city-mpg            205 non-null int64
24 highway-mpg         205 non-null int64
25 price               205 non-null object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

```
data.describe()
```

	symboling	wheel-base	length	width	height	curb-weight	engine-size	compression-ratio	city-mpg	highway-mpg
<b>count</b>	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
<b>mean</b>	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	10.142537	25.219512	30.751220
<b>std</b>	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	3.972040	6.542142	6.886443
<b>min</b>	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	7.000000	13.000000	16.000000
<b>25%</b>	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	8.600000	19.000000	25.000000
<b>50%</b>	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	9.000000	24.000000	30.000000
<b>75%</b>	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	9.400000	30.000000	34.000000
<b>max</b>	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	23.000000	49.000000	54.000000

## ▼ Cleaning Data

```
data['normalized-losses'].unique()
```

```

array(['?', '164', '158', '192', '188', '121', '98', '81', '118', '148',
       '110', '145', '137', '101', '78', '106', '85', '107', '104', '113',

```

```
'150', '129', '115', '93', '142', '161', '153', '125', '128',
'122', '103', '168', '108', '194', '231', '119', '154', '74',
'186', '83', '102', '89', '87', '77', '91', '134', '65', '197',
'90', '94', '256', '95'], dtype=object)
```

```
data.isnull().sum()
```

```
symboling          0
normalized-losses  0
make              0
fuel-type          0
aspiration         0
num-of-doors       0
body-style         0
drive-wheels       0
engine-location    0
wheel-base        0
length            0
width             0
height            0
curb-weight        0
engine-type        0
num-of-cylinders   0
engine-size        0
fuel-system        0
bore              0
stroke            0
compression-ratio  0
horsepower         0
peak-rpm          0
city-mpg           0
highway-mpg        0
price             0
dtype: int64
```

```
data.notnull().sum()
```

```
symboling          205
normalized-losses  205
```

```
make                205
fuel-type           205
aspiration           205
num-of-doors        205
body-style           205
drive-wheels         205
engine-location      205
wheel-base          205
length              205
width                205
height               205
curb-weight          205
engine-type          205
num-of-cylinders     205
engine-size          205
fuel-system          205
bore                 205
stroke               205
compression-ratio    205
horsepower           205
peak-rpm             205
city-mpg             205
highway-mpg          205
price                205
dtype: int64
```

```
data.duplicated().sum()
```

```
0
```

```
cols = data.columns
cols
```

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
       'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
       'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
       'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
       'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
```

```
'highway-mpg', 'price'],
dtype='object')
```

```
for col in cols:
```

```
    print(col,data[col].nunique())
```

```
symboling 6
normalized-losses 52
make 22
fuel-type 2
aspiration 2
num-of-doors 3
body-style 5
drive-wheels 3
engine-location 2
wheel-base 53
length 75
width 44
height 49
curb-weight 171
engine-type 7
num-of-cylinders 7
engine-size 44
fuel-system 8
bore 39
stroke 37
compression-ratio 32
horsepower 60
peak-rpm 24
city-mpg 29
highway-mpg 30
price 187
```

```
for col in cols:
```

```
    print(col,data[col].unique())
```

```
3750 3495 3770 3740 3685 3900 3715 2910 1918 1944 2004 2145 2370 2328
2833 2921 2926 2365 2405 2403 1889 2017 1938 1951 2028 1971 2037 2008
3334 3303 3005 3306 3060 3071 3130 3030 3107 3430 3075 3353 3305 3485
```

```

2324 2302 3095 3296 3060 3071 3139 3020 3197 3430 3075 3252 3285 3485
3130 2818 2778 2756 2800 3366 2579 2460 2658 2695 2707 2758 2808 2847
2050 2120 2240 2190 2340 2510 2290 2455 2420 2650 1985 2040 2015 2280
3110 2081 2109 2275 2094 2122 2140 2169 2204 2265 2300 2540 2536 2551
2679 2714 2975 2326 2480 2414 2458 2976 3016 3131 3151 2261 2209 2264
2212 2319 2254 2221 2661 2563 2912 3034 2935 3042 3045 3157 2952 3049
3012 3217 3062]
engine-type ['dohc' 'ohcv' 'ohc' 'l' 'rotor' 'ohcf' 'dohcv']
num-of-cylinders ['four' 'six' 'five' 'three' 'twelve' 'two' 'eight']
engine-size [130 152 109 136 131 108 164 209 61 90 98 122 156 92 79 110 111 119
258 326 91 70 80 140 134 183 234 308 304 97 103 120 181 151 194 203
132 121 146 171 161 141 173 145]
fuel-system ['mpfi' '2bbl' 'mfi' '1bbl' 'spfi' '4bbl' 'idi' 'spdi']
bore ['3.47' '2.68' '3.19' '3.13' '3.5' '3.31' '3.62' '2.91' '3.03' '2.97'
'3.34' '3.6' '2.92' '3.15' '3.43' '3.63' '3.54' '3.08' '?' '3.39' '3.76'
'3.58' '3.46' '3.8' '3.78' '3.17' '3.35' '3.59' '2.99' '3.33' '3.7'
'3.61' '3.94' '3.74' '2.54' '3.05' '3.27' '3.24' '3.01']
stroke ['2.68' '3.47' '3.4' '2.8' '3.19' '3.39' '3.03' '3.11' '3.23' '3.46' '3.9'
'3.41' '3.07' '3.58' '4.17' '2.76' '3.15' '?' '3.16' '3.64' '3.1' '3.35'
'3.12' '3.86' '3.29' '3.27' '3.52' '2.19' '3.21' '2.9' '2.07' '2.36'
'2.64' '3.08' '3.5' '3.54' '2.87']
compression-ratio [ 9. 10. 8. 8.5 8.3 7. 8.8 9.5 9.6 9.41 9.4 7.6
9.2 10.1 9.1 8.1 11.5 8.6 22.7 22. 21.5 7.5 21.9 7.8
8.4 21. 8.7 9.31 9.3 7.7 22.5 23. ]
horsepower ['111' '154' '102' '115' '110' '140' '160' '101' '121' '182' '48' '70'
'68' '88' '145' '58' '76' '60' '86' '100' '78' '90' '176' '262' '135'
'84' '64' '120' '72' '123' '155' '184' '175' '116' '69' '55' '97' '152'
'200' '95' '142' '143' '207' '288' '?' '73' '82' '94' '62' '56' '112'
'92' '161' '156' '52' '85' '114' '162' '134' '106']
peak-rpm ['5000' '5500' '5800' '4250' '5400' '5100' '4800' '6000' '4750' '4650'
'4200' '4350' '4500' '5200' '4150' '5600' '5900' '5750' '?' '5250' '4900'
'4400' '6600' '5300']
city-mpg [21 19 24 18 17 16 23 20 15 47 38 37 31 49 30 27 25 13 26 36 22 14 45 28
32 35 34 29 33]
highway-mpg [27 26 30 22 25 20 29 28 53 43 41 38 24 54 42 34 33 31 19 17 23 32 39 18
16 37 50 36 47 46]
price ['13495' '16500' '13950' '17450' '15250' '17710' '18920' '23875' '?'
'16430' '16925' '20970' '21105' '24565' '30760' '41315' '36880' '5151'
'6295' '6575' '5572' '6377' '7957' '6229' '6692' '7609' '8558' '8921'
'12964' '6479' '6855' '5399' '6529' '7129' '7295' '7895' '9095' '8845'
'10295' '12945' '10345' '6785' '11048' '32250' '35550' '36000' '5195'
'6005' '6705' '6605' '7305' '10045' '11045' '13045' '15045' '18405']

```

```

        6095      6795      6695      7595      10945      11845      13645      15045      8495
'10595' '10245' '10795' '11245' '18280' '18344' '25552' '28248' '28176'
'31600' '34184' '35056' '40960' '45400' '16503' '5389' '6189' '6669'
'7689' '9959' '8499' '12629' '14869' '14489' '6989' '8189' '9279' '5499'
'7099' '6649' '6849' '7349' '7299' '7799' '7499' '7999' '8249' '8949'
'9549' '13499' '14399' '17199' '19699' '18399' '11900' '13200' '12440'
'13860' '15580' '16900' '16695' '17075' '16630' '17950' '18150' '12764'
'22018' '32528' '34028' '37028' '9295' '9895' '11850' '12170' '15040'
'15510' '18620' '5118' '7053' '7603' '7126' '7775' '9960' '9233' '11259'
'7463' '10198' '8013' '11694' '5348' '6338' '6488' '6918' '7898' '8778'
'6938' '7198' '7788' '7738' '8358' '9258' '8058' '8238' '9298' '9538'
'8449' '9639' '9989' '11199' '11549' '17669' '8948' '10698' '9988'
'10898' '11248' '16558' '15998' '15690' '15750' '7975' '7995' '8195'
'9495' '9995' '11595' '9980' '13295' '13845' '12290' '12940' '13415'
'15985' '16515' '18420' '18950' '16845' '19045' '21485' '22470' '22625']

```

```
cl = ['normalized-losses', 'num-of-doors', 'bore', 'stroke', 'horsepower', 'peak-rpm', 'price']
```

```
for col in cl:
    print(data[col].value_counts())
```

```

100      2
176      2
55       1
262      1
134      1
115      1
140      1
48       1
58       1
60       1
78       1
135      1
200      1
64       1
120      1
72       1
154      1
288      1
143      1
142      1
175      1

```



```

175      1
106      1
Name: horsepower, dtype: int64
5500     37
4800     36
5000     27
5200     23
5400     13
6000      9
5250      7
4500      7
5800      7
4200      5
4150      5
4750      4
4350      4
5100      3
4250      3
5900      3
4400      3
?         2
6600      2
4650      1
5600      1
5750      1
4900      1
5300      1
Name: peak-rpm, dtype: int64
?         4
8921      2
18150     2
8845      2
8495      2
..
45400     1
16503     1
5389      1
6189      1
22625     1

```

```
for col in cl :
```

```
data[col].replace({'?': np.nan}, inplace = True)
```

```
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	

5 rows × 26 columns



```
cl = ['normalized-losses', 'num-of-doors', 'bore', 'stroke', 'horsepower', 'peak-rpm', 'price']
```

```
for col in cl:
    print(data[col].unique())
```

```
[nan '164' '158' '192' '188' '121' '98' '81' '118' '148' '110' '145' '137'
'101' '78' '106' '85' '107' '104' '113' '150' '129' '115' '93' '142'
'161' '153' '125' '128' '122' '103' '168' '108' '194' '231' '119' '154'
'74' '186' '83' '102' '89' '87' '77' '91' '134' '65' '197' '90' '94'
'256' '95']
['two' 'four' nan]
['3.47' '2.68' '3.19' '3.13' '3.5' '3.31' '3.62' '2.91' '3.03' '2.97'
'3.34' '3.6' '2.92' '3.15' '3.43' '3.63' '3.54' '3.08' nan '3.39' '3.76']
```

```

'3.58' '3.46' '3.8' '3.78' '3.17' '3.35' '3.59' '2.99' '3.33' '3.7'
'3.61' '3.94' '3.74' '2.54' '3.05' '3.27' '3.24' '3.01']
['2.68' '3.47' '3.4' '2.8' '3.19' '3.39' '3.03' '3.11' '3.23' '3.46' '3.9'
'3.41' '3.07' '3.58' '4.17' '2.76' '3.15' nan '3.16' '3.64' '3.1' '3.35'
'3.12' '3.86' '3.29' '3.27' '3.52' '2.19' '3.21' '2.9' '2.07' '2.36'
'2.64' '3.08' '3.5' '3.54' '2.87']
['111' '154' '102' '115' '110' '140' '160' '101' '121' '182' '48' '70'
'68' '88' '145' '58' '76' '60' '86' '100' '78' '90' '176' '262' '135'
'84' '64' '120' '72' '123' '155' '184' '175' '116' '69' '55' '97' '152'
'200' '95' '142' '143' '207' '288' nan '73' '82' '94' '62' '56' '112'
'92' '161' '156' '52' '85' '114' '162' '134' '106']
['5000' '5500' '5800' '4250' '5400' '5100' '4800' '6000' '4750' '4650'
'4200' '4350' '4500' '5200' '4150' '5600' '5900' '5750' nan '5250' '4900'
'4400' '6600' '5300']
['13495' '16500' '13950' '17450' '15250' '17710' '18920' '23875' nan
'16430' '16925' '20970' '21105' '24565' '30760' '41315' '36880' '5151'
'6295' '6575' '5572' '6377' '7957' '6229' '6692' '7609' '8558' '8921'
'12964' '6479' '6855' '5399' '6529' '7129' '7295' '7895' '9095' '8845'
'10295' '12945' '10345' '6785' '11048' '32250' '35550' '36000' '5195'
'6095' '6795' '6695' '7395' '10945' '11845' '13645' '15645' '8495'
'10595' '10245' '10795' '11245' '18280' '18344' '25552' '28248' '28176'
'31600' '34184' '35056' '40960' '45400' '16503' '5389' '6189' '6669'
'7689' '9959' '8499' '12629' '14869' '14489' '6989' '8189' '9279' '5499'
'7099' '6649' '6849' '7349' '7299' '7799' '7499' '7999' '8249' '8949'
'9549' '13499' '14399' '17199' '19699' '18399' '11900' '13200' '12440'
'13860' '15580' '16900' '16695' '17075' '16630' '17950' '18150' '12764'
'22018' '32528' '34028' '37028' '9295' '9895' '11850' '12170' '15040'
'15510' '18620' '5118' '7053' '7603' '7126' '7775' '9960' '9233' '11259'
'7463' '10198' '8013' '11694' '5348' '6338' '6488' '6918' '7898' '8778'
'6938' '7198' '7788' '7738' '8358' '9258' '8058' '8238' '9298' '9538'
'8449' '9639' '9989' '11199' '11549' '17669' '8948' '10698' '9988'
'10898' '11248' '16558' '15998' '15690' '15750' '7975' '7995' '8195'
'9495' '9995' '11595' '9980' '13295' '13845' '12290' '12940' '13415'
'15985' '16515' '18420' '18950' '16845' '19045' '21485' '22470' '22625']

```

```
data.isnull().sum()
```

```

symboling      0
normalized-losses  41
make           0

```

```

fuel-type      0
aspiration     0
num-of-doors   2
body-style     0
drive-wheels   0
engine-location 0
wheel-base     0
length         0
width          0
height         0
curb-weight    0
engine-type     0
num-of-cylinders 0
engine-size    0
fuel-system    0
bore           4
stroke         4
compression-ratio 0
horsepower     2
peak-rpm       2
city-mpg       0
highway-mpg    0
price          4
dtype: int64

```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null   int64
1   normalized-losses      164 non-null   object
2   make                   205 non-null   object
3   fuel-type              205 non-null   object
4   aspiration              205 non-null   object
5   num-of-doors           203 non-null   object
6   body-style             205 non-null   object
7   drive-wheels           205 non-null   object

```

```

8  engine-location    205 non-null    object
9  wheel-base        205 non-null    float64
10 length            205 non-null    float64
11 width             205 non-null    float64
12 height            205 non-null    float64
13 curb-weight       205 non-null    int64
14 engine-type       205 non-null    object
15 num-of-cylinders  205 non-null    object
16 engine-size       205 non-null    int64
17 fuel-system       205 non-null    object
18 bore              201 non-null    object
19 stroke            201 non-null    object
20 compression-ratio 205 non-null    float64
21 horsepower        203 non-null    object
22 peak-rpm          203 non-null    object
23 city-mpg          205 non-null    int64
24 highway-mpg       205 non-null    int64
25 price            201 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

```

string = ['normalized-losses','horsepower','peak-rpm', 'price','bore', 'stroke']
#data['normalized-losses'] = data['normalized-losses'].astype(float)
# data['horsepower'] = data['horsepower'].astype(float)
# data['peak-rpm'] = data['peak-rpm'].astype(float)
# data['price'] = data['price'].astype(float)
# data['bore'] = data['bore'].astype(float)
# data['stroke'] = data['stroke'].astype(float)

```

```

for s in string:
    data[s] = data[s].astype(float)

```

```

for col in cols:
    print(col,data[col].unique())

```

```

32.3 33. 33.1 33.3 33.8 36.2 37.3]
curb-weight [2548 2823 2337 2824 2507 2844 2954 3086 3053 2395 2710 2765 3055 3230
3380 3505 1488 1874 1909 1876 2128 1967 1989 2191 2535 2811 1713 1819
1837 1940 1956 2010 2024 2236 2289 2304 2372 2465 2293 2734 4066 3950
1800 1800 1805 1815 1850 2200 2205 2500 2410 2410 2405 2670 2700 2515

```

```

1890 1900 1905 1945 1950 2380 2385 2500 2410 2443 2425 2670 2700 3515
3750 3495 3770 3740 3685 3900 3715 2910 1918 1944 2004 2145 2370 2328
2833 2921 2926 2365 2405 2403 1889 2017 1938 1951 2028 1971 2037 2008
2324 2302 3095 3296 3060 3071 3139 3020 3197 3430 3075 3252 3285 3485
3130 2818 2778 2756 2800 3366 2579 2460 2658 2695 2707 2758 2808 2847
2050 2120 2240 2190 2340 2510 2290 2455 2420 2650 1985 2040 2015 2280

3110 2081 2109 2275 2094 2122 2140 2169 2204 2265 2300 2540 2536 2551
2679 2714 2975 2326 2480 2414 2458 2976 3016 3131 3151 2261 2209 2264
2212 2319 2254 2221 2661 2563 2912 3034 2935 3042 3045 3157 2952 3049
3012 3217 3062]
engine-type ['dohc' 'ohcv' 'ohc' 'l' 'rotor' 'ohcf' 'dohcv']
num-of-cylinders ['four' 'six' 'five' 'three' 'twelve' 'two' 'eight']
engine-size [130 152 109 136 131 108 164 209 61 90 98 122 156 92 79 110 111 119
258 326 91 70 80 140 134 183 234 308 304 97 103 120 181 151 194 203
132 121 146 171 161 141 173 145]
fuel-system ['mpfi' '2bbl' 'mfi' '1bbl' 'spfi' '4bbl' 'idi' 'spdi']
bore [3.47 2.68 3.19 3.13 3.5 3.31 3.62 2.91 3.03 2.97 3.34 3.6 2.92 3.15
3.43 3.63 3.54 3.08 nan 3.39 3.76 3.58 3.46 3.8 3.78 3.17 3.35 3.59
2.99 3.33 3.7 3.61 3.94 3.74 2.54 3.05 3.27 3.24 3.01]
stroke [2.68 3.47 3.4 2.8 3.19 3.39 3.03 3.11 3.23 3.46 3.9 3.41 3.07 3.58
4.17 2.76 3.15 nan 3.16 3.64 3.1 3.35 3.12 3.86 3.29 3.27 3.52 2.19
3.21 2.9 2.07 2.36 2.64 3.08 3.5 3.54 2.87]
compression-ratio [ 9. 10. 8. 8.5 8.3 7. 8.8 9.5 9.6 9.41 9.4 7.6
9.2 10.1 9.1 8.1 11.5 8.6 22.7 22. 21.5 7.5 21.9 7.8
8.4 21. 8.7 9.31 9.3 7.7 22.5 23. ]
horsepower [111. 154. 102. 115. 110. 140. 160. 101. 121. 182. 48. 70. 68. 88.
145. 58. 76. 60. 86. 100. 78. 90. 176. 262. 135. 84. 64. 120.
72. 123. 155. 184. 175. 116. 69. 55. 97. 152. 200. 95. 142. 143.
207. 288. nan 73. 82. 94. 62. 56. 112. 92. 161. 156. 52. 85.
114. 162. 134. 106.]
peak-rpm [5000. 5500. 5800. 4250. 5400. 5100. 4800. 6000. 4750. 4650. 4200. 4350.
4500. 5200. 4150. 5600. 5900. 5750. nan 5250. 4900. 4400. 6600. 5300.]
city-mpg [21 19 24 18 17 16 23 20 15 47 38 37 31 49 30 27 25 13 26 36 22 14 45 28
32 35 34 29 33]
highway-mpg [27 26 30 22 25 20 29 28 53 43 41 38 24 54 42 34 33 31 19 17 23 32 39 18
16 37 50 36 47 46]
price [13495. 16500. 13950. 17450. 15250. 17710. 18920. 23875. nan 16430.
16925. 20970. 21105. 24565. 30760. 41315. 36880. 5151. 6295. 6575.
5572. 6377. 7957. 6229. 6692. 7609. 8558. 8921. 12964. 6479.
6855. 5399. 6529. 7129. 7295. 7895. 9095. 8845. 10295. 12945.
10345. 6785. 11048. 32250. 35550. 36000. 5195. 6095. 6795. 6695.
7305. 10015. 11015. 12015. 15015. 20105. 10505. 10015. 10705. 11015

```

```

7395. 10945. 11845. 13645. 15645. 8495. 10595. 10245. 10795. 11245.
18280. 18344. 25552. 28248. 28176. 31600. 34184. 35056. 40960. 45400.
16503. 5389. 6189. 6669. 7689. 9959. 8499. 12629. 14869. 14489.
6989. 8189. 9279. 5499. 7099. 6649. 6849. 7349. 7299. 7799.
7499. 7999. 8249. 8949. 9549. 13499. 14399. 17199. 19699. 18399.
11900. 13200. 12440. 13860. 15580. 16900. 16695. 17075. 16630. 17950.
18150. 12764. 22018. 32528. 34028. 37028. 9295. 9895. 11850. 12170.
15040. 15510. 18620. 5118. 7053. 7603. 7126. 7775. 9960. 9233.
11259. 7463. 10198. 8013. 11694. 5348. 6338. 6488. 6918. 7898.
8778. 6938. 7198. 7788. 7738. 8358. 9258. 8058. 8238. 9298.
9538. 8449. 9639. 9989. 11199. 11549. 17669. 8948. 10698. 9988.
10898. 11248. 16558. 15998. 15690. 15750. 7975. 7995. 8195. 9495.
9995. 11595. 9980. 13295. 13845. 12290. 12940. 13415. 15985. 16515.
18470 18950 16845 19045 21485 22470 22625 1

```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	symboling	205 non-null	int64
1	normalized-losses	164 non-null	float64
2	make	205 non-null	object
3	fuel-type	205 non-null	object
4	aspiration	205 non-null	object
5	num-of-doors	203 non-null	object
6	body-style	205 non-null	object
7	drive-wheels	205 non-null	object
8	engine-location	205 non-null	object
9	wheel-base	205 non-null	float64
10	length	205 non-null	float64
11	width	205 non-null	float64
12	height	205 non-null	float64
13	curb-weight	205 non-null	int64
14	engine-type	205 non-null	object
15	num-of-cylinders	205 non-null	object
16	engine-size	205 non-null	int64
17	fuel-system	205 non-null	object
18	bore	201 non-null	float64

```
19  stroke                201 non-null    float64
20  compression-ratio     205 non-null    float64
21  horsepower             203 non-null    float64
22  peak-rpm               203 non-null    float64
23  city-mpg               205 non-null    int64
24  highway-mpg            205 non-null    int64
25  price                  201 non-null    float64
dtypes: float64(11), int64(5), object(10)
memory usage: 41.8+ KB
```

```
data['normalized-losses'] = data['normalized-losses'].fillna(data['normalized-losses'].mean())
data['horsepower'] = data['horsepower'].fillna(data['horsepower'].mean())
data['peak-rpm'] = data['peak-rpm'].fillna(data['peak-rpm'].mean())
data['price'] = data['price'].fillna(data['price'].mean())
data['bore'] = data['bore'].fillna(data['bore'].mean())
data['stroke'] = data['stroke'].fillna(data['stroke'].mean())
#Replacing '4wd' with 'fwd' in 'drivewheel' column
data['drive-wheels'] = data['drive-wheels'].replace('4wd','fwd')
```

```
data.isnull().sum()
```

```
symboling           0
normalized-losses   0
make                0
fuel-type           0
aspiration          0
num-of-doors        2
body-style          0
drive-wheels        0
engine-location     0
wheel-base         0
length             0
width              0
height             0
curb-weight         0
engine-type         0
num-of-cylinders    0
engine-size         0
```



```
fuel-system      0
bore             0
stroke          0
compression-ratio 0
horsepower      0
peak-rpm        0
city-mpg        0
highway-mpg     0
price           0
dtype: int64
```

```
data['num-of-doors'].value_counts()
```

```
four    114
two      89
Name: num-of-doors, dtype: int64
```

```
data['num-of-doors'].replace({np.nan : 'two'},inplace = True)
```

```
data['drive-wheels'].unique()
```

```
array(['rwd', 'fwd'], dtype=object)
```

```
data.isnull().sum()
```

```
symboling      0
normalized-losses 0
make          0
fuel-type      0
aspiration     0
num-of-doors   0
body-style     0
drive-wheels   0
engine-location 0
wheel-base    0
length        0
width         0
```

```

height          0
curb-weight     0
engine-type     0
num-of-cylinders 0
engine-size     0
fuel-system     0
bore            0
stroke          0
compression-ratio 0
horsepower      0
peak-rpm        0
city-mpg        0
highway-mpg     0
price           0
dtype: int64

```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null   int64
1   normalized-losses      205 non-null   float64
2   make                   205 non-null   object
3   fuel-type              205 non-null   object
4   aspiration              205 non-null   object
5   num-of-doors           205 non-null   object
6   body-style             205 non-null   object
7   drive-wheels           205 non-null   object
8   engine-location        205 non-null   object
9   wheel-base             205 non-null   float64
10  length                 205 non-null   float64
11  width                  205 non-null   float64
12  height                 205 non-null   float64
13  curb-weight            205 non-null   int64
14  engine-type            205 non-null   object
15  num-of-cylinders       205 non-null   object
16  engine-size            205 non-null   int64

```

```

17 fuel-system      205 non-null    object
18 bore            205 non-null    float64
19 stroke          205 non-null    float64
20 compression-ratio 205 non-null    float64
21 horsepower      205 non-null    float64
22 peak-rpm        205 non-null    float64
23 city-mpg        205 non-null    int64
24 highway-mpg     205 non-null    int64
25 price           205 non-null    float64

```

```
dtypes: float64(11), int64(5), object(10)
```

```
memory usage: 41.8+ KB
```

```
data.describe().style.background_gradient()
```

	symboling	normalized- losses	wheel- base	length	width	height	curb- weight	engine- size	bore	stroke	cor
<b>count</b>	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	
<b>mean</b>	0.834146	122.000000	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.329751	3.255423	
<b>std</b>	1.245307	31.681008	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.270844	0.313597	
<b>min</b>	-2.000000	65.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.540000	2.070000	
<b>25%</b>	0.000000	101.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.150000	3.110000	
<b>50%</b>	1.000000	122.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.310000	3.290000	
<b>75%</b>	2.000000	137.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.580000	3.410000	
<b>max</b>	3.000000	256.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	3.940000	4.170000	

```
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	122.0	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	1	122.0	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
3	2	164.0	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	
4	2	164.0	audi	gas	std	four	sedan	fwd	front	99.4	...	136	mpfi	3.19	

5 rows × 26 columns

## ▼ Handelling Outliers

- We can see that there are (205-88)= 117 records, which are outliers in the dataset.

```
# Finding outliers in all the numerical columns with 1.5 IQR rule and removing the outlier records
# #cul = ['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
#        'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
#        'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
#        'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
#        'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
#        'highway-mpg', 'price']

col_numeric=['wheel-base', 'normalized-losses', 'length', 'width', 'height', 'curb-weight', 'engine-size', 'bore', 'stroke',
             'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price']

for col in col_numeric:
    q1 = data[col].quantile(0.25)
    q3 = data[col].quantile(0.75)
```

```
iqr = q3-q1
range_low = q1-1.5*iqr
range_high = q3+1.5*iqr
data = data.loc[(data[col] > range_low) & (data[col] < range_high)]
```

```
data.shape
```

```
(121, 26)
```

## ▼ Checking Data Imbalance


- We can see that there is data imbalance in below columns:-
  - symboling - There are very few with rating -2.
  - fuletype - All the cars fule type is Gas, as Diesel cars were removed while removing outliers..
  - aspiration - Lesser number of turbo than std.
  - engineloaction - All the engine location is in front, as all the rear engine cars were removed while removing outliers.
  - enginetype - Considerably more number of ohc than others.
  - cylindernumber - Large number of four cyliners than others.
  - fulesystem - mpfi and 2bbl fulesystem cars are more comparitavely others.
  - CarCompany - Most of the Toyota company cars were surveyed.


```
# Listing categorical columns for checking data imbalance and plotting them
```

```
# #cul = ['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
#        'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
#        'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
#        'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
#        'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
#        'highway-mpg', 'price']
```

```
col_category = ['symboling', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels', 'engine-location', 'engine-type',
                'num-of-cylinders', 'fuel-system', 'make']
```

```
k=0
plt.figure(figsize=(20,25))
for col in col_category:
    k=k+1
    plt.subplot(4, 3,k)
    data[col].value_counts().plot(kind='bar');
    plt.title(col)
```






## ▼ Visualising the data to check the possibility of linear regression model


- We can see that there are few columns that have linear relationship with the target variable "price". So, we can build a linear regression model here.

```
# Visualising the numerical variables
plt.figure(figsize=(12,12))
sns.pairplot(data[col_numeric])
plt.show()
```



&lt;Figure size 864x864 with 0 Axes&gt;



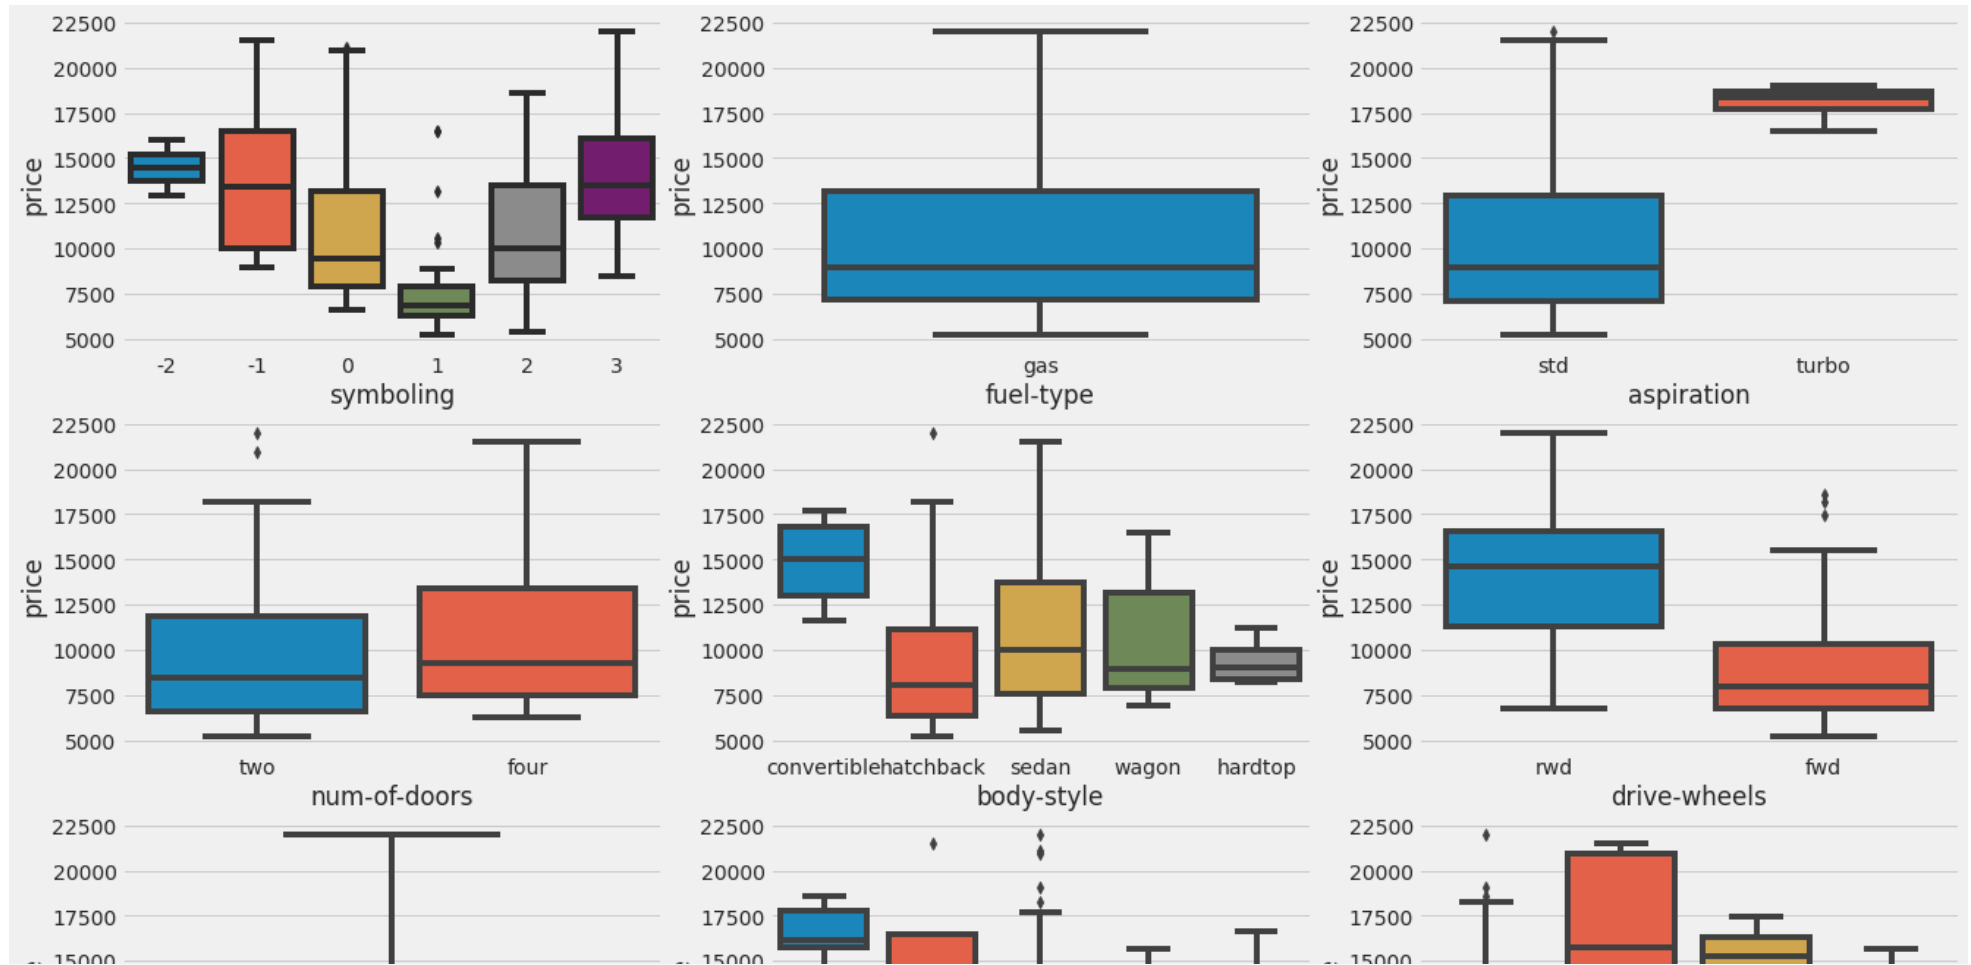


## ▼ Visualising the categorical variables


- CarCompany - Porsche has very high median price compared to other cars, though the number of Porsche cars is very less. Volvo, alfa-romero, audi and BMW are also high median price than others. Saab has wide range of price, with high median price.
  - aspiration - std has lower median than turbo.
  - carbody - convertible has higher median than others.
  - symboling - -2 and -1 have higher median price than others.
  - enginelocation - rear has very high median price than front.
  - cylindernumber - Four has lower median than others.
  - fulesystem - 1bbl and 2bbl have lower median price than others.
- Now at least we know that what are the variables have impact on the price. So as which variables are important for the model building.

```
# Boxplot for all categorical variables except CarCompany
```

```
# As X labels are not clearly visible for CarCompany. It is plotted in the next cell with bigger figure size.  
k=0  
plt.figure(figsize=(20,18))  
for col in range (len(col_category)-1):  
    k=k+1  
    plt.subplot(4, 3, k)  
    ax = sns.boxplot(x = col_category[col], y = 'price', data = data)
```



```
plt.figure(figsize=(15,8))
ax = sns.boxplot(x = 'make', y = 'price', data = data)
temp = ax.set_xticklabels(ax.get_xticklabels(), rotation = 45, horizontalalignment='right')
```



## ▼ Step 2:- Preparing the data for model building

### Encoding

- Converting categorical variables (fueltype, aspiration, doornumber, drivewheel, enginelocation) with two levels to binary variables.

```
data['fuel-type'].unique()  
  
array(['gas'], dtype=object)
```

```
# fueltype
# Convert "gas" to 1 and "diesel" to 0
data['fuel-type'] = data['fuel-type'].map({'gas': 1, 'diesel': 0})
data.head()
```

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	...	engine- size	fuel- system	bore	sti
<b>0</b>	3	122.0	alfa- romero	1	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
<b>1</b>	3	122.0	alfa- romero	1	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
<b>2</b>	1	122.0	alfa- romero	1	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
<b>3</b>	2	164.0	audi	1	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	
<b>4</b>	2	164.0	audi	1	std	four	sedan	fwd	front	99.4	...	136	mpfi	3.19	

5 rows × 26 columns



```
# aspiration
# Convert "std" to 1 and "turbo" to 0
data['aspiration'] = data['aspiration'].map({'std':1, 'turbo':0})
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	122.0	alfa-romero	1	1	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	122.0	alfa-romero	1	1	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	1	122.0	alfa-romero	1	1	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
3	2	164.0	audi	1	1	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	
4	2	164.0	audi	1	1	four	sedan	fwd	front	99.4	...	136	mpfi	3.19	

```
# #cul = ['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
#         'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
#         'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
#         'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
#         'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
#         'highway-mpg', 'price']
```

```
data['num-of-doors'].unique()
```

```
array(['two', 'four'], dtype=object)
```

```
# doornumber
# Convert "four" to 1 and "two" to 0
data['num-of-doors'] = data['num-of-doors'].map({'four':1, 'two':0})
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	122.0	alfa-romero	1	1	0	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	122.0	alfa-romero	1	1	0	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	1	122.0	alfa-romero	1	1	0	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
3	2	164.0	audi	1	1	1	sedan	fwd	front	99.8	...	109	mpfi	3.19	
4	2	164.0	audi	1	1	1	sedan	fwd	front	99.4	...	136	mpfi	3.19	

5 rows × 26 columns

```
data['drive-wheels'].unique()
```

```
array(['rwd', 'fwd'], dtype=object)
```

```
# drivewheel
# Convert "fwd" to 1 and "rwd" to 0
data['drive-wheels'] = data['drive-wheels'].map({'fwd':1, 'rwd':0})
data.head()
```



	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	122.0	alfa-romero	1	1	0	convertible	0	front	88.6	...	130	mpfi	3.47	
1	3	122.0	alfa-romero	1	1	0	convertible	0	front	88.6	...	130	mpfi	3.47	

```
data['engine-location'].unique()
```

```
array(['front'], dtype=object)
```

```
4      2      164.0    audi      1      1      1      sedan      1      front      99.4    ...      136      mpfi      3.19
```

```
# enginelocation
```

```
# Convert "front" to 1 and "rear" to 0
```

```
data['engine-location'] = data['engine-location'].map({'front':1, 'rear':0})
```

```
data.head()
```

symboling normalized-losses make fuel-tvne aspiration num-of- body-style drive-wheels engine-location wheel-base ... engine-size fuel-svstem bore sti


## ▼ Dummy variables

- Converting other categorical variables with more than two levels to dummy variables. We have to create (n-1) dummy variables by removing the base status. n is the number of levels of the variables.

```
# Creating dummy variables for 'symboling'
# Dropping the redundant dummy variable (-2)
symboling_status = pd.get_dummies(data['symboling'],drop_first=True)
symboling_status.head()
```

	-1	0	1	2	3
0	0	0	0	0	1
1	0	0	0	0	1
2	0	0	1	0	0
3	0	0	0	1	0
4	0	0	0	1	0

```
# Renaming column names for better readability
symboling_status = symboling_status.rename(columns={-1:'symboling(-1)', 0:'symboling(0)', 1:'symboling(1)',2:'symboling(2)', 3:'symboling(3)'})
symboling_status.head()
```

	symboling(-1)	symboling(0)	symboling(1)	symboling(2)	symboling(3)	
<b>0</b>	0	0	0	0	1	
<b>1</b>	0	0	0	0	1	

```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,symboling_status], axis=1)
data.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	horsepower	peak-rpm	city-mpg
<b>0</b>	3	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	...	111.0	5000.0	21
<b>1</b>	3	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	...	111.0	5000.0	21
<b>2</b>	1	122.0	alfa-romero	1	1	0	hatchback	0	1	94.5	...	154.0	5000.0	19
<b>3</b>	2	164.0	audi	1	1	1	sedan	1	1	99.8	...	102.0	5500.0	24
<b>4</b>	2	164.0	audi	1	1	1	sedan	1	1	99.4	...	115.0	5500.0	18

5 rows × 31 columns



```
# Dropping the 'symboling' column as we don't need it anymore
data = data.drop('symboling',axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	length	...	horsepower	peak- rpm	city- mpg	hi
0	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	168.8	...	111.0	5000.0	21	
1	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	168.8	...	111.0	5000.0	21	
2	122.0	alfa-romero	1	1	0	hatchback	0	1	94.5	171.2	...	154.0	5000.0	19	
3	164.0	audi	1	1	1	sedan	1	1	99.8	176.6	...	102.0	5500.0	24	
4	164.0	audi	1	1	1	sedan	1	1	99.4	176.6	...	115.0	5500.0	18	

5 rows × 30 columns



```
# Creating dummy variables for 'carbody'
# Dropping the redundant dummy variable (convertible)
carbody_status = pd.get_dummies(data['body-style'],drop_first=True)
carbody_status.head()
```

	hardtop	hatchback	sedan	wagon
0	0	0	0	0
1	0	0	0	0
2	0	1	0	0
3	0	0	1	0
4	0	0	1	0



```
# Renaming column names for better readability
carbody_status = carbody_status.rename(columns={'hardtop':'carbody(hardtop)', 'hatchback':'carbody(hatchback)', 'sedan':'carbody(sedan)', 'wagon':'carbody(wagon)'})
carbody_status.head()
```

```
carbody_status.head()
```

	carbody(hardtop)	carbody(hatchback)	carbody(sedan)	carbody(wagon)	
0	0	0	0	0	
1	0	0	0	0	
2	0	1	0	0	
3	0	0	1	0	
4	0	0	1	0	

```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,carbody_status], axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	length	...	price	symboling(-1)	symbol
0	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	168.8	...	13495.0	0	
1	122.0	alfa-romero	1	1	0	convertible	0	1	88.6	168.8	...	16500.0	0	
2	122.0	alfa-romero	1	1	0	hatchback	0	1	94.5	171.2	...	16500.0	0	
3	164.0	audi	1	1	1	sedan	1	1	99.8	176.6	...	13950.0	0	
4	164.0	audi	1	1	1	sedan	1	1	99.4	176.6	...	17450.0	0	

5 rows × 34 columns



```
# Dropping the 'symboling' column as we don't need it
data = data.drop('body-style',axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	price	symboling(-1)	symboling
0	122.0	alfa- romero	1	1	0	0	1	88.6	168.8	64.1	...	13495.0	0	
1	122.0	alfa- romero	1	1	0	0	1	88.6	168.8	64.1	...	16500.0	0	
2	122.0	alfa- romero	1	1	0	0	1	94.5	171.2	65.5	...	16500.0	0	
3	164.0	audi	1	1	1	1	1	99.8	176.6	66.2	...	13950.0	0	
4	164.0	audi	1	1	1	1	1	99.4	176.6	66.4	...	17450.0	0	

5 rows × 33 columns



```
cul = ['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base',
'length', 'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower',
'peak-rpm', 'city-mpg', 'highway-mpg', 'price']
```

```
# Creating dummy variables for 'engine-type'
# Dropping the redundant dummy variable (dohc)
engine_type_status = pd.get_dummies(data['engine-type'], drop_first=True)
engine_type_status.head()
```

	1	ohc	ohcv	rotor
0	0	0	0	0
1	0	0	0	0
2	0	0	1	0
3	0	1	0	0
4	0	1	0	0



```
# Renaming column name for better readability
enginetype_status = enginetype_status.rename(columns={'dohcv':'enginetype(dohcv)', 'l':'enginetype(l)', 'ohc':'enginetype(ohc)',
                                                    'ohcf':'enginetype(ohcf)', 'ohcv':'enginetype(ohcv)', 'rotor':'enginetype(rotor)'})
enginetype_status.head()
```

	enginetype(l)	enginetype(ohc)	enginetype(ohcv)	enginetype(rotor)
0	0	0	0	0
1	0	0	0	0
2	0	0	1	0
3	0	1	0	0
4	0	1	0	0



```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,enginetype_status], axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	symboling(2)	symboling(3)	carb
<b>0</b>	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	1	
<b>1</b>	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	1	
<b>2</b>	122.0	alfa-romero	1	1	0	0	1	94.5	171.2	65.5	...	0	0	
<b>3</b>	164.0	audi	1	1	1	1	1	99.8	176.6	66.2	...	1	0	

```
# Dropping the 'enginetype' column as we don't need it
data = data.drop('engine-type',axis=1)
data.head()
```



```
# Creating dummy variables for 'cylindernumber'
# Dropping the redundant dummy variable (eight)
cylindernumber_status = pd.get_dummies(data['num-of-cylinders'], drop_first=True)
cylindernumber_status.head()
```

	four	six	two
0	1	0	0
1	1	0	0
2	0	1	0
3	1	0	0
4	0	0	0

```
# Renaming column name for better readability
cylindernumber_status = cylindernumber_status.rename(columns={'five':'cylindernumber(five)', 'four':'cylindernumber(four)', 'six':'cylindernumber(six)', 'three':'cylindernumber(three)', 'twelve':'cylindernumber(twelve)', 'two':'cylindernumber(two)'})
cylindernumber_status.head()
```

	cylindernumber(four)	cylindernumber(six)	cylindernumber(two)
0	1	0	0
1	1	0	0
2	0	1	0
3	1	0	0
4	0	0	0

```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,cylindernumber_status], axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	carbody(hatchback)	carbody(seda
<b>0</b>	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
<b>1</b>	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
<b>2</b>	122.0	alfa-romero	1	1	0	0	1	94.5	171.2	65.5	...	1	
<b>3</b>	164.0	audi	1	1	1	1	1	99.8	176.6	66.2	...	0	
<b>4</b>	164.0	audi	1	1	1	1	1	99.4	176.6	66.4	...	0	


5 rows × 39 columns



```
# Dropping the 'cylindernumber' column as we don't need it
data = data.drop('num-of-cylinders',axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	carbody(hatchback)	carbody(seda
0	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
1	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	

```
# Creating dummy variables for 'fuelsystem'
# Dropping the redundant dummy variable (1bb1)
fuelsystem_status = pd.get_dummies(data['fuel-system'], drop_first=True)
fuelsystem_status.head()
```

	2bb1	4bb1	mpfi	spfi	
0	0	0	1	0	
1	0	0	1	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	0	1	0	

```
# Renaming column name for better readability
fuelsystem_status = fuelsystem_status.rename(columns={'2bb1':'fuelsystem(2bb1)', '4bb1':'fuelsystem(4bb1)', 'idi':'fuelsystem(idi)',
'mfi':'fuelsystem(mfi)', 'mpfi':'fuelsystem(mpfi)', 'spdi':'fuelsystem(spdi)',
'spfi':'fuelsystem(spfi)'})
fuelsystem_status.head()
```

	fuelsystem(2bbl)	fuelsystem(4bbl)	fuelsystem(mpfi)	fuelsystem(spfi)	
0	0	0	1	0	
1	0	0	1	0	

```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,fuelsystem_status], axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	enginetype(ohc)	enginetype(ohcv)
0	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
1	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
2	122.0	alfa-romero	1	1	0	0	1	94.5	171.2	65.5	...	0	
3	164.0	audi	1	1	1	1	1	99.8	176.6	66.2	...	1	
4	164.0	audi	1	1	1	1	1	99.4	176.6	66.4	...	1	

5 rows × 42 columns



```
# Dropping the 'fuelsystem' column as we don't need it
data = data.drop('fuel-system',axis=1)
data.head()
```

	normalized- losses	make	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	...	engine type(ohc)	engine type(ohcv)
0	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
1	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
2	122.0	alfa-romero	1	1	0	0	1	94.5	171.2	65.5	...	0	
3	164.0	audi	1	1	1	1	1	99.8	176.6	66.2	...	1	
4	164.0	audi	1	1	1	1	1	99.4	176.6	66.4	...	1	

5 rows × 41 columns



```
# Creating dummy variables for 'CarCompany'
# Dropping the redundant dummy variable (alfa-romero)
CarCompany_status = pd.get_dummies(data['make'], drop_first=True)
CarCompany_status.head()
```

	audi	bmw	chevrolet	dodge	honda	isuzu	mazda	mercury	mitsubishi	nissan	peugot	plymouth	porsche	saab	toyota	volks
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```
# Renaming column name for better readability
CarCompany_status = CarCompany_status.rename(columns={'audi':'CarCompany(audi)', 'bmw':'CarCompany(bmw)', 'buick':'CarCompany(buick)'})
```

```
'chevrolet': 'CarCompany(chevrolet)', 'dodge': 'CarCompany(dodge)', 'honda': 'CarCo
'isuzu': 'CarCompany(isuzu)', 'jaguar': 'CarCompany(jaguar)', 'mazda': 'CarCompany(m
'mercury': 'CarCompany(mercury)', 'mitsubishi': 'CarCompany(mitsubishi)', 'nissan':
'peugeot': 'CarCompany(peugeot)', 'plymouth': 'CarCompany(plymouth)', 'porsche': 'Ca
'renault': 'CarCompany(renault)', 'saab': 'CarCompany(saab)', 'subaru': 'CarCompany(
'toyota': 'CarCompany(toyota)', 'volkswagen': 'CarCompany(volkswagen)', 'volvo': 'Ca
```

```
CarCompany_status.head()
```

	CarCompany(audi)	CarCompany(bmw)	CarCompany(chevrolet)	CarCompany(dodge)	CarCompany(honda)	CarCompany(isuzu)	CarCompan
0	0	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	0	0	0	0	
3	1	0	0	0	0	0	
4	1	0	0	0	0	0	



```
# Concating the dummy dataframe with original dataframe
data = pd.concat([data,CarCompany_status], axis=1)
data.head()
```

	normalized-losses	make	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	width	...	CarCompany(mercury)	CarCompany(
0	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
1	122.0	alfa-romero	1	1	0	0	1	88.6	168.8	64.1	...	0	
		alfa											

```
# Dropping the 'CarCompany' column as we don't need it
data = data.drop('make',axis=1)
data.head()
```

	normalized-losses	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	width	height	...	CarCompany(mercury)	CarCompany(
0	122.0	1	1	0	0	1	88.6	168.8	64.1	48.8	...	0	
1	122.0	1	1	0	0	1	88.6	168.8	64.1	48.8	...	0	
2	122.0	1	1	0	0	1	94.5	171.2	65.5	52.4	...	0	
3	164.0	1	1	1	1	1	99.8	176.6	66.2	54.3	...	0	
4	164.0	1	1	1	1	1	99.4	176.6	66.4	54.3	...	0	

5 rows × 57 columns



```
x = data.loc[:,['normalized-losses', 'fuel-type', 'aspiration', 'num-of-doors',
'drive-wheels', 'engine-location', 'wheel-base', 'length', 'width',
'height', 'curb-weight', 'engine-size', 'bore', 'stroke',
'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
'highway-mpg','symboling(-1)', 'symboling(0)', 'symboling(1)'],
```

```
'symboling(2)', 'symboling(3)', 'carbody(hardtop)',
'carbody(hatchback)', 'carbody(sedan)', 'carbody(wagon)',
'enginetype(1)', 'enginetype(ohc)', 'enginetype(ohcv)',
'enginetype(rotor)', 'cylindernumber(four)', 'cylindernumber(six)',
'cylindernumber(two)', 'fuelsystem(2bbl)', 'fuelsystem(4bbl)',
'fuelsystem(mpfi)', 'fuelsystem(spfi)', 'CarCompany(audi)',
'CarCompany(bmw)', 'CarCompany(chevrolet)', 'CarCompany(dodge)',
'CarCompany(honda)', 'CarCompany(isuzu)', 'CarCompany(mazda)',
'CarCompany(mercury)', 'CarCompany(mitsubishi)', 'CarCompany(nissan)',
'peugot', 'CarCompany(plymouth)', 'CarCompany(porsche)',
'CarCompany(saab)', 'CarCompany(toyota)', 'CarCompany(volkswagen)',
'CarCompany(volvo)']]
```

```
y = data.loc[:, 'price']
```

```
x.head()
```

	normalized- losses	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width	height	...	CarCompany(mercury)	CarCompany(
<b>0</b>	122.0	1	1	0	0	1	88.6	168.8	64.1	48.8	...	0	
<b>1</b>	122.0	1	1	0	0	1	88.6	168.8	64.1	48.8	...	0	
<b>2</b>	122.0	1	1	0	0	1	94.5	171.2	65.5	52.4	...	0	
<b>3</b>	164.0	1	1	1	1	1	99.8	176.6	66.2	54.3	...	0	
<b>4</b>	164.0	1	1	1	1	1	99.4	176.6	66.4	54.3	...	0	

5 rows × 56 columns





```
y.head()
```

```
0    13495.0
1    16500.0
2    16500.0
3    13950.0
4    17450.0
Name: price, dtype: float64
```

```
data.info()
```

```
1  fuel-type          121 non-null  int64
2  aspiration         121 non-null  int64
3  num-of-doors       121 non-null  int64
4  drive-wheels       121 non-null  int64
5  engine-location    121 non-null  int64
6  wheel-base         121 non-null  float64
7  length            121 non-null  float64
8  width             121 non-null  float64
9  height            121 non-null  float64
10 curb-weight       121 non-null  int64
11 engine-size       121 non-null  int64
12 bore             121 non-null  float64
13 stroke           121 non-null  float64
14 compression-ratio 121 non-null  float64
15 horsepower        121 non-null  float64
16 peak-rpm          121 non-null  float64
17 city-mpg          121 non-null  int64
18 highway-mpg       121 non-null  int64
19 price             121 non-null  float64
20 symboling(-1)     121 non-null  uint8
21 symboling(0)      121 non-null  uint8
22 symboling(1)      121 non-null  uint8
23 symboling(2)      121 non-null  uint8
24 symboling(3)      121 non-null  uint8
25 carbody(hardtop)   121 non-null  uint8
26 carbody(hatchback) 121 non-null  uint8
27 carbody(sedan)    121 non-null  uint8
28 carbody(wagon)    121 non-null  uint8
29 enginetype(1)     121 non-null  uint8
30 enginetype(ohc)   121 non-null  uint8
31 enginetype(stroke) 121 non-null  uint8
```

```

31 enginetype(oncev)      121 non-null  uint8
32 enginetype(rotor)      121 non-null  uint8
33 cylindernumber(four)   121 non-null  uint8
34 cylindernumber(six)    121 non-null  uint8
35 cylindernumber(two)    121 non-null  uint8
36 fuelsystem(2bbl)       121 non-null  uint8
37 fuelsystem(4bbl)       121 non-null  uint8
38 fuelsystem(mpfi)       121 non-null  uint8
39 fuelsystem(spfi)       121 non-null  uint8
40 CarCompany(audi)       121 non-null  uint8
41 CarCompany(bmw)        121 non-null  uint8
42 CarCompany(chevrolet)  121 non-null  uint8
43 CarCompany(dodge)      121 non-null  uint8
44 CarCompany(honda)      121 non-null  uint8
45 CarCompany(isuzu)      121 non-null  uint8
46 CarCompany(mazda)      121 non-null  uint8
47 CarCompany(mercury)    121 non-null  uint8
48 CarCompany(mitsubishi) 121 non-null  uint8
49 CarCompany(nissan)     121 non-null  uint8
50 peugot                 121 non-null  uint8
51 CarCompany(plymouth)   121 non-null  uint8
52 CarCompany(porsche)    121 non-null  uint8
53 CarCompany(saab)       121 non-null  uint8
54 CarCompany(toyota)     121 non-null  uint8
55 CarCompany(volkswagen) 121 non-null  uint8
56 CarCompany(volvo)      121 non-null  uint8
dtypes: float64(11), int64(9), uint8(37)
memory usage: 28.3 KB

```

## ▼ Splitting Data into Train and test

```

#Importing Model Building Libraries

# Train-test split
from sklearn.model_selection import train_test_split
# Min-max sciling
from sklearn.preprocessing import MinMaxScaler

```

```
# Statsmodel
import statsmodels.api as sm
# VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor
#R-squared
from sklearn.metrics import r2_score
# Label encoding
from sklearn.preprocessing import LabelEncoder
# Importing RFE
from sklearn.feature_selection import RFE
# Importing LinearRegression
from sklearn.linear_model import LinearRegression
# Supress warning
import warnings
warnings.filterwarnings('ignore')
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=42)
```

```
from sklearn import metrics
from sklearn.model_selection import cross_val_score

def cross_val(model):
    pred = cross_val_score(model, x, y, cv=10)
    return pred.mean()

def print_evaluate(true, predicted):
    mae = metrics.mean_absolute_error(true, predicted)
    mse = metrics.mean_squared_error(true, predicted)
    rmse = np.sqrt(metrics.mean_squared_error(true, predicted))
    r2_square = metrics.r2_score(true, predicted)
    print('MAE:', mae)
    print('MSE:', mse)
    print('RMSE:', rmse)
    print('R2 Square', r2_square)
    print('_____')

def evaluate(true, predicted):
```

```
def evaluate(true, predicted):  
    mae = metrics.mean_absolute_error(true, predicted)  
    mse = metrics.mean_squared_error(true, predicted)  
    rmse = np.sqrt(metrics.mean_squared_error(true, predicted))  
    r2_square = metrics.r2_score(true, predicted)  
    return mae, mse, rmse, r2_square
```

```
scaler = MinMaxScaler()  
  
x_train = scaler.fit_transform(x_train)  
  
x_test = scaler.transform(x_test)
```

## ▼ Model Creation Linear Regression


```
lin_reg = LinearRegression(normalize=True)  
lin_reg.fit(x_train,y_train)
```

```
LinearRegression(normalize=True)
```

```
# print the intercept  
print(lin_reg.intercept_)
```

```
12401.400923571686
```

```
coeff_df = pd.DataFrame(lin_reg.coef_, x.columns, columns=['Coefficient'])  
coeff_df
```

	Coefficient 
<b>normalized-losses</b>	5.364262e+02
<b>fuel-type</b>	4.911271e-11
<b>aspiration</b>	-7.135967e+02
<b>num-of-doors</b>	4.517465e+02
<b>drive-wheels</b>	-1.911465e+02
<b>engine-location</b>	-9.822543e-11
<b>wheel-base</b>	1.214284e+04
<b>length</b>	-4.232021e+03
<b>width</b>	2.269482e+03
<b>height</b>	-3.810101e+03
<b>curb-weight</b>	8.231102e+03
<b>engine-size</b>	-1.809028e+04
<b>bore</b>	1.584240e+04
<b>stroke</b>	4.138864e+03
<b>compression-ratio</b>	-6.888009e+02
<b>horsepower</b>	1.711234e+03
<b>peak-rpm</b>	1.221798e+04
<b>city-mpg</b>	-3.038019e+03
<b>highway-mpg</b>	5.841129e+03
<b>symboling(-1)</b>	6.174041e+02
<b>symboling(0)</b>	-2.006351e+03
<b>symboling(1)</b>	-1.546794e+03

<b>symboling(2)</b>	-1.557762e+03
<b>symboling(3)</b>	-1.043796e+03
<b>carbody(hardtop)</b>	-5.198632e+03
<b>carbody(hatchback)</b>	-3.961535e+03
<b>carbody(sedan)</b>	-3.077945e+03
<b>carbody(wagon)</b>	-3.113475e+03
<b>enginetype(l)</b>	-7.024955e+03
<b>enginetype(ohc)</b>	-1.134631e+03
<b>enginetype(ohcv)</b>	7.816698e+03
<b>enginetype(rotor)</b>	-6.064400e+03
<b>cylindernumber(four)</b>	-5.229821e+03
<b>cylindernumber(six)</b>	-3.517498e+03
<b>cylindernumber(two)</b>	-6.064400e+03
<b>fuelsystem(2bbl)</b>	1.171292e+03
<b>fuelsystem(4bbl)</b>	-6.064400e+03
<b>fuelsystem(mpfi)</b>	5.702351e+02
<b>fuelsystem(spfi)</b>	-3.637979e-11
<b>CarCompany(audi)</b>	-7.266831e+03
<b>CarCompany(bmw)</b>	1.257492e+04
<b>CarCompany(chevrolet)</b>	-1.097124e+04
<b>CarCompany(dodge)</b>	-1.218712e+04
<b>CarCompany(honda)</b>	-1.120816e+04
<b>CarCompany(isuzu)</b>	-5.549611e+03

<b>CarCompany(mazda)</b>	-6.300397e+03
<b>CarCompany(mercury)</b>	-7.275958e-12
<b>CarCompany(mitsubishi)</b>	-1.119795e+04
<b>CarCompany(nissan)</b>	-9.381603e+03
<b>peugot</b>	-7.024955e+03
<b>CarCompany(plymouth)</b>	-1.172562e+04
<b>CarCompany(porsche)</b>	-3.389551e+03
<b>CarCompany(saab)</b>	-7.219976e+03
<b>CarCompany(toyota)</b>	-5.800348e+03

## ▼ Prediction from our model

```
pred = lin_reg.predict(x_test)
pred
```

```
array([10327.23884945, 23890.93919519, 16804.43226042, 6594.76993172,
       10174.6998297 , 9426.91386408, 6884.37240061, 5043.34217433,
       18535.82106196, 8170.20742958, 7555.64493054, 6044.76139966,
       5080.18061249, 8262.77502563, 10495.33986194, 18173.83170366,
       9874.61397272, 16966.04542699, 8140.4676725 , 9896.05780829,
       14324.84419744, 9527.24163496, 12199.6243466 , 11430.09044116,
       10327.23884945])
```

```
test_pred = lin_reg.predict(x_test)
train_pred = lin_reg.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)
```

```
results_df = pd.DataFrame(data=[["Linear Regression", *evaluate(y_test, test_pred) , cross_val(LinearRegression())]],
                          columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
```

Test set evaluation:

---

MAE: 1631.7548865140263  
 MSE: 5568857.322725073  
 RMSE: 2359.8426478740216  
 R2 Square 0.5950985663635897

---

Train set evaluation:

---

MAE: 659.7588443098801  
 MSE: 671090.4995822677  
 RMSE: 819.2011349981565  
 R2 Square 0.9626935377493883

---

## ▼ Robust Regression

### Random Sample Consensus - RANSAC

```
from sklearn.linear_model import RANSACRegressor

model = RANSACRegressor(base_estimator=LinearRegression(), max_trials=100)
model.fit(x_train, y_train)

test_pred = model.predict(x_test)
train_pred = model.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
```



```
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Robust Regression", *evaluate(y_test, test_pred) , cross_val(RANSACRegressor())]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
results_df = results_df.append(results_df_2, ignore_index=True)
```

Test set evaluation:

---

MAE: 2850.377532020369  
 MSE: 21900968.845926296  
 RMSE: 4679.8470964259395  
 R2 Square -0.5923794002685214

=====

Train set evaluation:

---

MAE: 1297.6926676300573  
 MSE: 11395141.908167949  
 RMSE: 3375.6691052542383  
 R2 Square 0.366534868692002

---

## ▼ Ridge Regression

```
from sklearn.linear_model import Ridge

model = Ridge(alpha=100, solver='cholesky', tol=0.0001, random_state=42)
model.fit(x_train, y_train)
pred = model.predict(x_test)

test_pred = model.predict(x_test)
train_pred = model.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
```

```

print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Ridge Regression", *evaluate(y_test, test_pred) , cross_val(Ridge())]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
results_df = results_df.append(results_df_2, ignore_index=True)

```

Test set evaluation:

---

MAE: 1869.0625915053417  
MSE: 5144109.611156661  
RMSE: 2268.0629645485287  
R2 Square 0.6259811958477433

---

=====

Train set evaluation:

---

MAE: 2196.7318893822303  
MSE: 8107280.544626266  
RMSE: 2847.328668177642  
R2 Square 0.5493097342586634

---

## ▼ Lasso Regression

```

from sklearn.linear_model import Lasso

model = Lasso(alpha=0.1,
               precompute=True,
               #       warm_start=True,
               positive=True,
               selection='random',
               random_state=42)

model.fit(x_train, y_train)

```

```

test_pred = model.predict(x_test)
train_pred = model.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Lasso Regression", *evaluate(y_test, test_pred) , cross_val(Lasso())]],
                           columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
results_df = results_df.append(results_df_2, ignore_index=True)

```

Test set evaluation:

---

MAE: 1086.3260714321107  
MSE: 2140202.7459618445  
RMSE: 1462.9431793346741  
R2 Square 0.8443897715647555

---

=====

Train set evaluation:

---

MAE: 1072.5992746327288  
MSE: 2072005.5320534536  
RMSE: 1439.4462588278361  
R2 Square 0.8848155409550171

---

## ▼ Elastic Net

```

from sklearn.linear_model import ElasticNet

model = ElasticNet(alpha=0.1, l1_ratio=0.9, selection='random', random_state=42)

```

```

model.fit(x_train, y_train)

test_pred = model.predict(x_test)
train_pred = model.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Elastic Net Regression", *evaluate(y_test, test_pred) , cross_val(ElasticNet())]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
results_df = results_df.append(results_df_2, ignore_index=True)

```

Test set evaluation:

---

MAE: 1207.3732628964087  
MSE: 2418734.7579660695  
RMSE: 1555.228201250887  
R2 Square 0.8241382182498722

=====

Train set evaluation:

---

MAE: 1023.4487602922059  
MSE: 1802685.98602198  
RMSE: 1342.641421237249  
R2 Square 0.8997872317830489

---

## ▼ Polynomial Regression

```

from sklearn.preprocessing import PolynomialFeatures

```

```

poly_reg = PolynomialFeatures(degree=2)

```

```

poly_reg = PolynomialFeatures(degree=2)

x_train_2_d = poly_reg.fit_transform(x_train)
x_test_2_d = poly_reg.transform(x_test)

lin_reg = LinearRegression(normalize=True)
lin_reg.fit(x_train_2_d,y_train)

test_pred = lin_reg.predict(x_test_2_d)
train_pred = lin_reg.predict(x_train_2_d)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Polynomial Regression", *evaluate(y_test, test_pred), 0]],
                           columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
results_df = results_df.append(results_df_2, ignore_index=True)

```

Test set evaluation:

---

MAE: 5679253330676.771  
MSE: 7.540490466357226e+25  
RMSE: 8683599752612.522  
R2 Square -5.48255274504984e+18

---

=====

Train set evaluation:

---

MAE: 37.48772090070482  
MSE: 24549.048928726774  
RMSE: 156.68136114013936  
R2 Square 0.9986352985659639

---

## ▼ Stochastic Gradient Descent

```
from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor(n_iter_no_change=250, penalty=None, eta0=0.0001, max_iter=100000)
sgd_reg.fit(x_train, y_train)

test_pred = sgd_reg.predict(x_test)
train_pred = sgd_reg.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Stochastic Gradient Descent", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
results_df = results_df.append(results_df_2, ignore_index=True)
```

Test set evaluation:

---

MAE: 1101.647249790968  
MSE: 2469197.5754734827  
RMSE: 1571.3680585634554  
R2 Square 0.820469159056938

=====

Train set evaluation:

---

MAE: 1249.030230590947  
MSE: 2604907.2553506363  
RMSE: 1613.9725076192085  
R2 Square 0.8551911042570646

---

## ▼ Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor

rf_reg = RandomForestRegressor(n_estimators=1000)
rf_reg.fit(x_train, y_train)

test_pred = rf_reg.predict(x_test)
train_pred = rf_reg.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)

print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Random Forest Regressor", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
results_df = results_df.append(results_df_2, ignore_index=True)
```

Test set evaluation:

---

MAE: 862.9970688166312  
 MSE: 1519722.824088209  
 RMSE: 1232.7703857929948  
 R2 Square 0.889503732176392

Train set evaluation:

---

MAE: 478.478963731933  
 MSE: 431517.79972999886  
 RMSE: 656.900144413136  
 R2 Square 0.9760115774010882

---

## ▼ Support Vector Machine

```

from sklearn.svm import SVR

svm_reg = SVR(kernel='rbf', C=1000000, epsilon=0.001)
svm_reg.fit(x_train, y_train)

test_pred = svm_reg.predict(x_test)
train_pred = svm_reg.predict(x_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)

print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["SVM Regressor", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
results_df = results_df.append(results_df_2, ignore_index=True)

```

Test set evaluation:

---

MAE: 1349.8107789830328  
MSE: 3968916.7746342896  
RMSE: 1992.2140383589033  
R2 Square 0.7114273182264543

---

Train set evaluation:

---

MAE: 174.7889295273078  
MSE: 267005.22538857773  
RMSE: 516.7254835873471  
R2 Square 0.9851569641235043

---

results\_df